

```
In [7]: import tensorflow as tf
```

Task:2 The dataset

Import MNIST

```
In [10]: from tensorflow.keras.datasets import mnist
(x_train,y_train),(x_test,y_test)=mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11490434/11490434 [=====] - 3s 0us/step

shapes of Imported Arrays

```
In [12]: print('x_train shape',x_train.shape)
print('y_tains shape',y_train.shape)
print('x_test shape',x_test.shape)
print('y_test shape',y_test.shape)
```

x\_train shape (60000, 28, 28)

y\_tains shape (60000,)


x\_test shape (10000, 28, 28)

y\_test shape (10000,)

Plot an Image Example

```
In [18]: from matplotlib import pyplot as plt
%matplotlib inline

plt.imshow(x_train[0], cmap='binary')
plt.show()
```



Display labels

```
In [19]: y_train[0]
```

```
Out[19]: 5
```

```
In [20]: print(set(y_train))
```

```
{0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
```

One Hot Encoding

Encoding Labels

```
In [21]: from tensorflow.keras.utils import to_categorical
```

```
In [23]: y_train_encoded= to_categorical(y_train)
y_test_encoded= to_categorical(y_test)
```

Validated Shapes

```
In [24]: print('y_train_encoded shape:',y_train_encoded.shape)
print('y_test_encoded shape:', y_test_encoded.shape)
```

```
y_train_encoded shape: (60000, 10)
```

```
y_test_encoded shape: (10000, 10)
```

Display Encoded Labels

```
In [26]: y_train_encoded[0]
```

```
Out[26]: array([0., 0., 0., 0., 0., 1., 0., 0., 0., 0.], dtype=float32)
```

Task 4: Neural Networks

Linear Equations

Task5 : Preprocessing the Examples

Inrolling N-dimensional Arrays to Vectors

```
In [30]: import numpy as np
x_train_resaped=np.reshape(x_train, (60000,784))
x_test_resaped=np.reshape(x_test,(10000,784))
```

```
print('x_train_resaped',x_train_resaped)
print('x_test_resaped',x_test_resaped)
```

```
x_train_resaped [[0 0 0 ... 0 0 0]
```

```
[0 0 0 ... 0 0 0]
```

```
[0 0 0 ... 0 0 0]
```

```
...
```

```
[0 0 0 ... 0 0 0]
```

```
[0 0 0 ... 0 0 0]
```

```
[0 0 0 ... 0 0 0]]
```

```
x_train_resaped [[0 0 0 ... 0 0 0]
```

```
[0 0 0 ... 0 0 0]
```

```
[0 0 0 ... 0 0 0]
```

```
...
```

```
[0 0 0 ... 0 0 0]
```

```
[0 0 0 ... 0 0 0]
```

```
[0 0 0 ... 0 0 0]]
```

## Display Pixel Values

```
In [40]: print(set(x_test_reshaped[0]))
```

{0, 129, 1, 3, 133, 5, 9, 140, 14, 17, 18, 19, 21, 22, 151, 159, 31, 163, 36, 35, 166, 38, 40, 170, 44, 52, 182, 185, 58, 59, 60, 187, 62, 57, 61, 66, 67, 198, 72, 75, 203, 205, 77, 207, 209, 83, 84, 219, 221, 222, 224, 225, 121, 227, 126, 229, 233, 106, 236, 238, 240, 241, 114, 115, 242, 248, 249, 250, 251, 253, 254, 255}

## Data Normalization

```
In [43]: x_mean=np.mean(x_train_reshaped)
x_std=np.std(x_train_reshaped)
epsilon =1e-10

x_train_norm=(x_train_reshaped-x_mean)/(x_std+epsilon)
x_test_norm=(x_test_reshaped-x_mean)/(x_std + epsilon)
```

## Display Normalized Pixel Values

```
In [44]: print(set(x_train_norm[0]))
```

{-0.38589016215482896, 1.306921966983251, 1.17964285952926, 1.803310486053816, 1.6887592893452241, 2.8215433456857437, 2.719720059722551, 1.1923707702746593, 1.7396709323268205, 2.057868700961798, 2.3633385588513764, 2.096052433197995, 1.7651267538176187, 2.7960875241949457, 2.7451758812133495, 2.45243393406917, 0.02140298169794222, -0.22042732246464067, 1.2305545025108566, 0.2759611966059242, 2.210603629906587, 2.6560805059955555, 2.6051688630139593, -0.4240738943910262, 0.4668798577869107, 0.1486820891519332, 0.3905123933145161, 1.0905474843114664, -0.09314821501064967, 1.4851127174188385, 2.7579037919587486, 1.5360243604004349, 0.07231462467953861, -0.13133194724684696, 1.294194056237852, 0.03413089244334132, 1.3451056992194483, 2.274243183633583, -0.24588314395543887, 0.772349715676489, 0.75962180493109, 0.7214380726948927, 0.1995937321335296, -0.41134598364562713, 0.5687031437501034, 0.5941589652409017, 0.9378125553666773, 0.9505404661120763, 0.6068868759863008, 0.4159682148053143, -0.042236572029053274, 2.7706317027041476, 2.1342361654341926, 0.12322626766113501, -0.08042030426525057, 0.16140999989733232, 1.8924058612716097, 1.2560103240016547, 2.185147808415789, 0.6196147867316999, 1.943317504253206, -0.11860403650144787, -0.30952269768243434, 1.9942291472348024, -0.2840668761916362, 2.6306246845047574, 2.286971094378982, -0.19497150097384247, -0.39861807290022805, 0.2886891073513233, 1.7523988430722195, 2.3887943803421745, 2.681536327486354, 1.4596568959280403, 2.439706023323771, 2.7833596134495466, 2.490617666305367, -0.10587612575604877, 1.5614801818912332, 1.9051337720170087, 1.6123918248728295, 1.268738234747054, 1.9560454149986053, 2.6433525952501564, 1.026907930584471}

## Task6: Creatin a Model

### Creating the model

```
In [48]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

model= Sequential([

    Dense(128,activation='relu',input_shape=(784,)),
    Dense(128,activation='relu'),
    Dense(10,activation='softmax')

])
```

```
2023-12-13 17:27:54.618855: E external/local_xla/xla/stream_executor/cuda/cuda_driver.cc:274] failed call to cuInit: CUDA_ERROR_NO_DEVICE: no CUDA-capable device is detected
```

## Compiling the Model

```
In [53]: model.compile(  
    optimizer='sgd',  
    loss='categorical_crossentropy',  
    metrics=['accuracy']  
)  
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	100480
dense_1 (Dense)	(None, 128)	16512
dense_2 (Dense)	(None, 10)	1290

=====  
Total params: 118282 (462.04 KB)  
Trainable params: 118282 (462.04 KB)  
Non-trainable params: 0 (0.00 Byte)  
=====

## Task7: Training the Model

### Training the Model

```
In [54]: model.fit(x_train_norm,y_train_encoded,epochs=3)  
  
Epoch 1/3  
1875/1875 [=====] - 9s 4ms/step - loss: 0.3760 - accuracy: 0.8899  
Epoch 2/3  
1875/1875 [=====] - 7s 4ms/step - loss: 0.1853 - accuracy: 0.9459  
Epoch 3/3  
1875/1875 [=====] - 7s 4ms/step - loss: 0.1398 - accuracy: 0.9592
```

```
Out[54]: <keras.src.callbacks.History at 0x7fb54a2b7370>
```

## Evaluating the Model

```
In [57]: _, accuracy=model.evaluate(x_test_norm,y_test_encoded)  
print('Test set accuracy:',accuracy*100)  
  
313/313 [=====] - 1s 4ms/step - loss: 0.1359 - accuracy: 0.9585  
Test set accuracy: 95.85000276565552
```

## Task 8 : Predictions

### Prediction on Test Set

```
In [58]: preds=model.predict(x_test_norm)  
print('Shape of preds:',preds.shape)
```

313/313 [=====] - 1s 3ms/step  
Shape of preds: (10000, 10)

Plotting the result

```
In [59]: plt.figure(figsize=(12,12))
```

Out[59]: <Figure size 1200x1200 with 0 Axes>

<Figure size 1200x1200 with 0 Axes>


```
In [64]: start_index=0

for i in range(25):
    plt.subplot(5,5,i+1)
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])


    pred=np.argmax(preds[start_index])
    gt =y_test[start_index+i]

    col='g'
    if pred != gt:
        col='r'

    plt.xlabel('i={},pred={},gt={}'.format(start_index+i,pred,gt))
    plt.imshow(x_test[start_index+i],cmap='binary')
plt.show()
```



```
In [67]: plt.plot(preds[8])
plt.show()
```



In [ ]: